

BEGINNER STUDY GUIDE

Exploring the Depths of Large Language Models

Key Concepts in Generative AI

12 Core Concepts · Explained Simply

© 2026 David Grunwald. All rights reserved.

This material is provided for personal study and educational use.

Introduction

This guide unpacks the essential concepts every beginner should understand when studying Large Language Models (LLMs) and Generative AI. Each concept is defined in plain language, with a short explanation of why it matters. Use the summary table at the end for quick review.

1. Model Architecture & Representation

Tokenization

What it is: The process of breaking raw text into smaller units called **tokens** — which can be words, sub-words, or characters.

Before a model can process language, sentences must be chopped into these standardized pieces. For example, "unbelievable" might become "un", "believ", "able". Tokenization is the very first step in the pipeline that turns human text into something a model can work with.

Key fact: Tokenization is the process that transforms text into the units the model reads. (Not to be confused with embeddings, which come next.)

Embeddings

What it is: The process of transforming tokens into **numerical (vector) representations** that the model can mathematically understand.

Computers don't understand words — they understand numbers. Embeddings map each token to a list of numbers (a vector) so that words with similar meanings sit close together in mathematical space. This is how a model "understands" that *king* and *queen* are related.

Key fact: Embeddings convert text into numerical representations understandable by the LLM.

Parameters

What it is: The internal values (weights) a model learns during training. The parameter count is a rough measure of a model's size and capacity.

Increasing the number of parameters generally lets a model capture more complex patterns. The trade-off is that larger models need more storage, memory, and compute.

Key fact: Increasing parameters increases the complexity of the output, providing more detail.

2. Training & Fine-Tuning

Epoch

What it is: One complete pass of the model through the entire training dataset. The number of epochs = how many times the model is exposed to the data.

Too few epochs and the model underfits (hasn't learned enough); too many and it can overfit (memorizes instead of generalizing). Tuning epochs is a balancing act.

Key fact: The **epoch** parameter represents the number of times a model is exposed to the training data.

Learning Rate

What it is: A setting that controls how big a step the model takes when adjusting its parameters during training.

A high learning rate trains fast but risks overshooting the best solution; a low learning rate is more precise but slower. As the learning task grows more complex, you often **lower the learning rate** so the model can converge carefully.

Key fact: Depending on the complexity of the learning rate, you lower the learning rate for more stable, careful training.

Fine-Tuning

What it is: Taking a pre-trained model and further training it on a specialized dataset so it performs better on a specific task.

Instead of building a model from scratch, you adapt a general model to your needs. Common use cases include **sentiment analysis** (judging if text is positive/negative) and **language translation**.

Key fact: Two common fine-tuning use cases: sentiment analysis and language translation.

LoRA — Low-Rank Adaptation

What it is: An efficient fine-tuning technique. **LoRA** stands for **Low-Rank Adaptation**.

Rather than updating every parameter (expensive), LoRA freezes the original model and trains a small number of extra low-rank matrices. This makes fine-tuning far cheaper and faster while keeping most of the original model intact.

Key fact: LoRA = Low-Rank Adaptation — a lightweight fine-tuning method.

3. Generation & Capabilities

Controllable Text Generation

What it is: Techniques that let a user guide or constrain the model's output (tone, style, format, topic) rather than accepting whatever it produces.

This is valuable when you need output tailored to a purpose. Its two main advantages: it's **helpful when customization is required**, and it **gives the user more flexibility**.

Key fact: Advantages — helpful if customization is required, and provides more flexibility.

Multimodal Generative AI

What it is: AI that works across multiple data types at once — text, images, audio, video, and sensor data.

By combining modes, these systems handle richer real-world tasks. Two common use cases are **computer vision** (interpreting images/video) and **robotics** (combining perception and action).

Key fact: Two common use cases: computer vision and robotics.

4. Limitations & Risks

General-Purpose LLM Limitations

What it is: A general-purpose LLM is trained broadly rather than for one specific job — which brings trade-offs.

Two key disadvantages: it **can't determine whether an event is "true"** (it predicts plausible text, not verified fact), and it **cannot understand context beyond its training data** (its knowledge is frozen at training time).

Key fact: Disadvantages — can't determine if an event is "true"; cannot understand context beyond training.

Bias in Machine Learning

What it is: When a model consistently produces skewed results because of imbalances in its training data or assumptions.

Bias doesn't cause random mistakes — it causes **systemic prediction error**, meaning the model is wrong in a consistent, patterned direction. This is a fairness and reliability concern.

Key fact: Bias causes **systemic prediction error**.

Current Limitations of LLMs

What it is: Broad challenges facing today's large models.

Three notable limitations:

- **Bias** of models (skewed, systematic outputs)
- **Accuracy** of models (they can produce confident but wrong answers)
- **Capacity** for storing these models (huge size demands significant storage/compute)

Key fact: Three limitations — bias, accuracy, and storage capacity.

Quick-Review Summary

Concept	One-Line Takeaway
Tokenization	Breaks text into tokens — the first processing step
Embeddings	Turns text into numerical vectors the LLM understands
Parameters	More parameters → more complex, detailed output
Epoch	Number of times the model sees the full training data
Learning Rate	Step size in training; lower it for complex tasks
Fine-Tuning	Adapt a pre-trained model (e.g., sentiment analysis, translation)
LoRA	Low-Rank Adaptation — efficient fine-tuning
Controllable Generation	Customization + flexibility over output
Multimodal AI	Works across data types (e.g., computer vision, robotics)
General-Purpose Limits	Can't verify truth; no context beyond training
Bias	Causes systemic (patterned) prediction error
LLM Limitations	Bias, accuracy, storage capacity